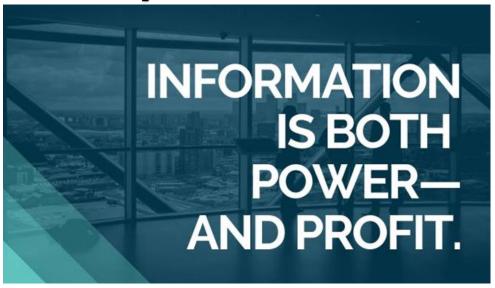


### Learning from data:

## Technical Nuances of Machine Learning in Plant Breeding

Alencar Xavier, 02/04/2021
Research Scientist at Corteva Biostatistics
Adjunct professor at Purdue University

## Adequate use of





### **Outline**

Three faces on machine learning

- 1. Good
- 2. Bad
- 3. Ugly





#### 1. Good

- Intro and motivation
- Filters in PB

#### 2. Bad

- Metrics of success
- Adequate validation

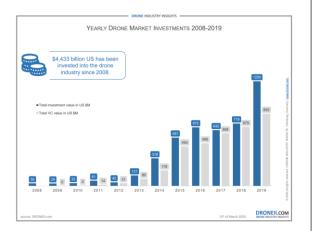
#### 3. Ugly

- Optimization
- From RR to NN

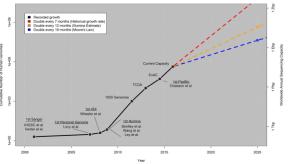
### Part 1 – Good learners



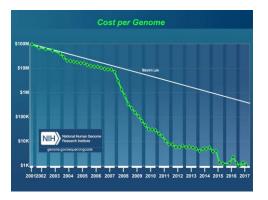
### More Pheno



# More Geno Growth of DNA Sequencing



The Cost of Sequencing a Human Genome. NIH. https://www.genome.gov/27565109/the-cost-of-sequencing-a-human-genome/



Stephens, Z. D.et al. (2015). Big data: astronomical or genomical? *PLoS biology*, *13*(7), e1002195.

## More Env \*\*\*

- UC Merced GridMET
- NWS NOAA
- NASA GISS
- Harmonized SoilDB
- USDA SSURGO

# More Computing 2

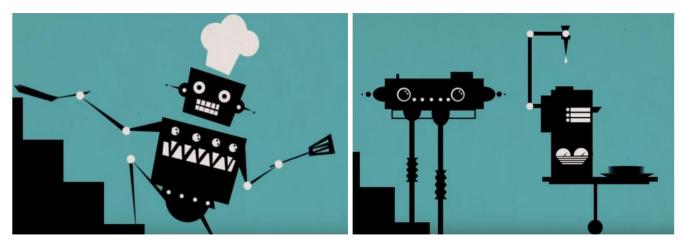






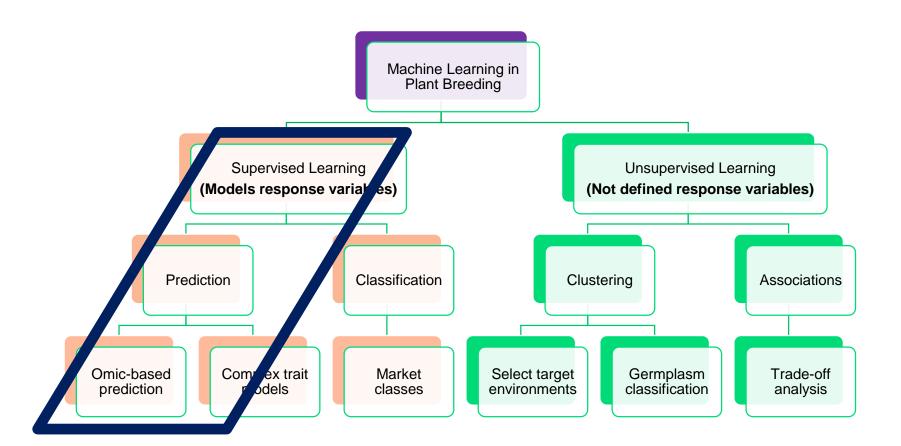
### Why is machine learning good for?

#### Good for solving single well-defined problem



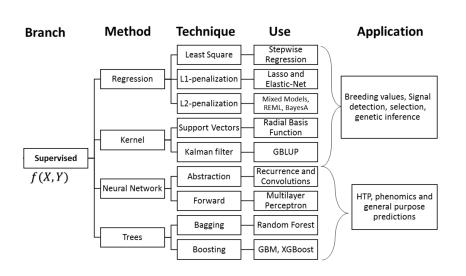
Source: https://www.youtube.com/watch?v=MPR3o6Hnf2g

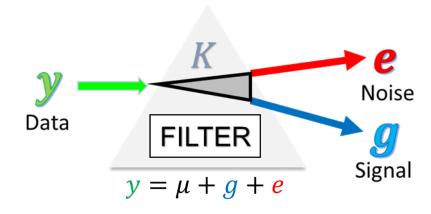






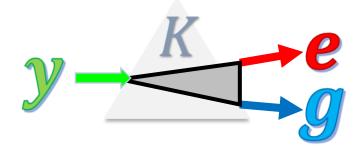
### Key of supervised learning: FILTERING

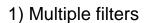


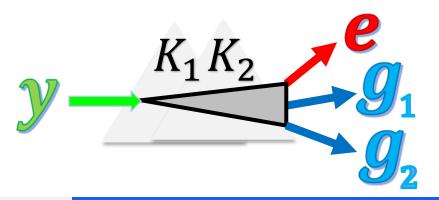




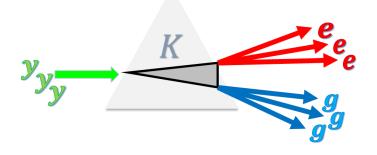
#### **Generalizations of simple filters**





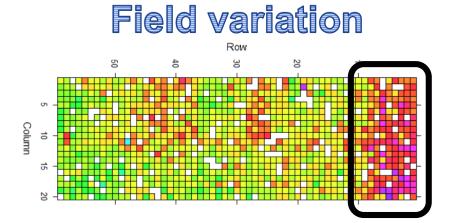


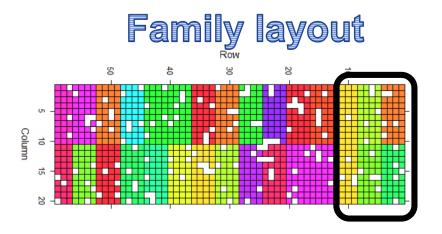






### 1) When bother with multiple filters?



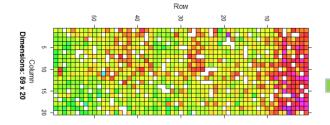


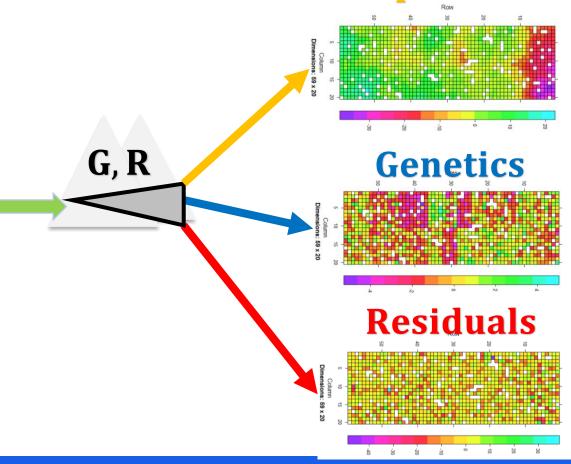
Some families were placed on unfavorable side of the field...

SoyNAM field, Indiana 2014



### **Pheno**



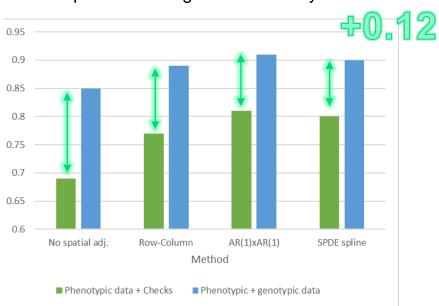


**Spatial** 

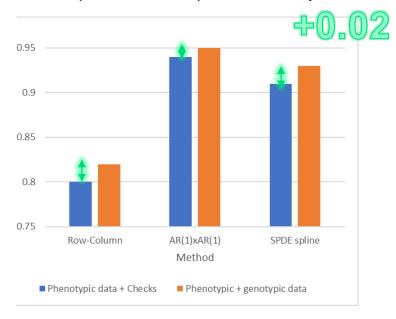


### Multiple filters benefit separability of signals

#### Improvement in genetic accuracy



#### Improvement in spatial accuracy

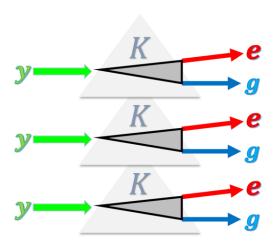


Results derived from simulation of field with 1500 plots, 1275 non-replicated entries + checks, from single trials

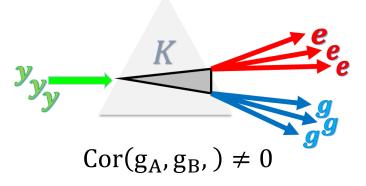


### 2) When bother with multi-response filters?

Filter one input at a time (parallelizable)



Filter passing multiple input at once



#### **New information**

Genetic correlation table

,	TICLIC COITCIALIOIT LAD			
		$g_1$	$g_2$	$g_3$
	$g_1$	1	$ ho_{12}$	$ ho_{13}$
	$g_2$	$ ho_{21}$	1	$\rho_{23}$
	$g_3$	$ ho_{31}$	$\rho_{32}$	1





#### Bivariate model

#### **Objective gain of information!**

$$\begin{bmatrix} y_1 \\ y_2 \end{bmatrix} = \begin{bmatrix} g_1 \\ g_2 \end{bmatrix} + \begin{bmatrix} e_1 \\ e_2 \end{bmatrix}, \quad y \sim N(0, G \otimes \Sigma_a + I \otimes \Sigma_e)$$

Model equation

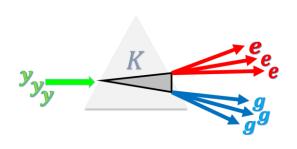
$$\begin{bmatrix} Z_1' \Sigma_e^{11} Z_1 + G^{-1} \Sigma_a^{11} & Z_1' \Sigma_e^{12} Z_2 + G^{-1} \Sigma_a^{12} \\ Z_2' \Sigma_e^{12} Z_1 + G^{-1} \Sigma_a^{12} & Z_2' \Sigma_e^{22} Z_2 + G^{-1} \Sigma_a^{22} \end{bmatrix} \begin{bmatrix} g_1 \\ g_2 \end{bmatrix} = \begin{bmatrix} Z_1' (\Sigma_e^{11} y_1 + \Sigma_e^{12} y_2) \\ Z_2' (\Sigma_e^{22} y_2 + \Sigma_e^{12} y_1) \end{bmatrix}$$

· Univariate vs bivariate solution for a given predictor

$$\begin{split} \mathbf{g}_1 &= (\mathbf{Z}_1' \boldsymbol{\Sigma}_e^{11} \mathbf{Z}_1 + \mathbf{G}^{-1} \boldsymbol{\Sigma}_a^{11})^{-1} (\mathbf{Z}_1' \boldsymbol{\Sigma}_e^{11} \mathbf{y}_1) \\ \\ \mathbf{g}_1 | \mathbf{g}_2 &= (\mathbf{Z}_1' \boldsymbol{\Sigma}_e^{11} \mathbf{Z}_1 + \mathbf{G}^{-1} \boldsymbol{\Sigma}_a^{11})^{-1} (\mathbf{Z}_1' (\boldsymbol{\Sigma}_e^{11} \mathbf{y}_1 + \boldsymbol{\Sigma}_e^{12} \mathbf{y}_2) - (\mathbf{Z}_1' \boldsymbol{\Sigma}_e^{12} \mathbf{Z}_2 + \mathbf{G}^{-1} \boldsymbol{\Sigma}_a^{12}) \mathbf{g}_2) \end{split}$$

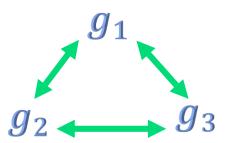


### **Sparse & Directed Filtering**



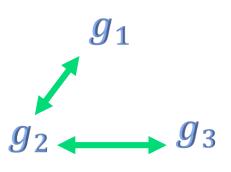
#### Genetic correlation table

	$g_1$	$g_2$	$g_3$
$g_1$	1	$ ho_{12}$	$ ho_{13}$
$g_2$	$ ho_{21}$	1	$ ho_{23}$
$g_3$	$ ho_{31}$	$ ho_{32}$	1



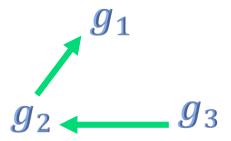
### **Sparse**

	$g_1$	$g_2$	$g_3$
$g_1$	1	$ ho_{12}$	l
$g_2$	$ ho_{21}$	1	$\rho_{23}$
$g_3$	_	$\rho_{32}$	1



#### **Directed**

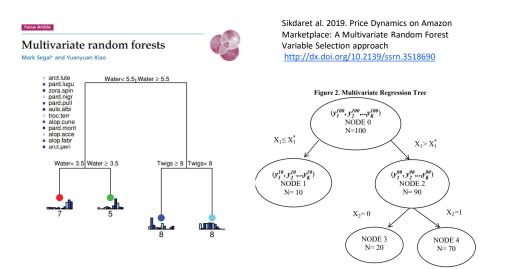
	$g_1$	$g_2$	$g_3$
$g_1$	1	ı	ı
$g_2$	$ ho_{21}$	1	_
$g_3$	_	$ ho_{32}$	1



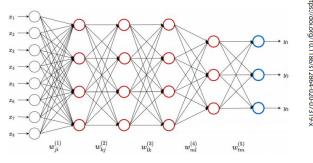


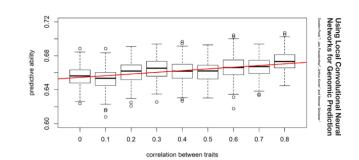
### Modern multi-response machinery

MVRF feature splits occur for all responses together



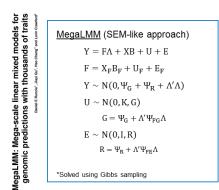
#### Multi-response comes natural for DNNs

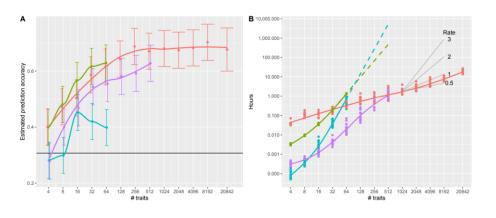






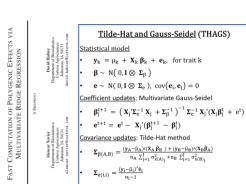
### Modern multi-response machinery

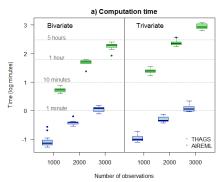


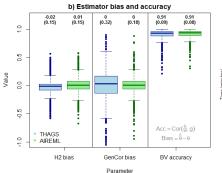


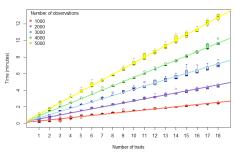
→ MegaLMM → MCMCglmm → MTG2 → phenix

"Classical" mixed model framework has been evolving in this area









#### 1. Good

- Intro and motivation
- Filters in PB

#### 2. Bad

- Metrics of success
- Adequate validation

#### 3. Ugly

- Optimization
- From RR to NN

(avoiding)

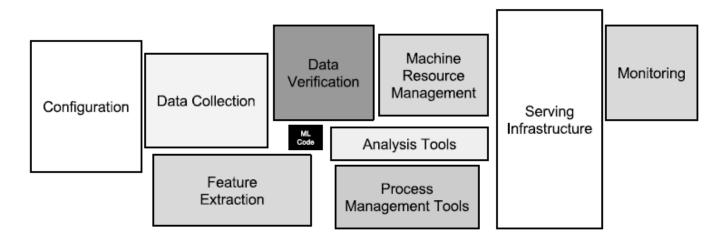
### Part 2 – Bad learning



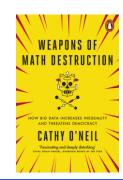
### There is usually more to ML than a proof of concept with cross-validations

#### **Hidden Technical Debt in Machine Learning Systems**

doi/10.5555/2969442.2969519



- How easily can an entirely new algorithmic approach be tested at full scale?
- What is the transitive closure of all data dependencies?
- How precisely can the impact of a new change to the system be measured?
- Does improving one model or signal degrade others?
- How quickly can new members of the team be brought up to speed?

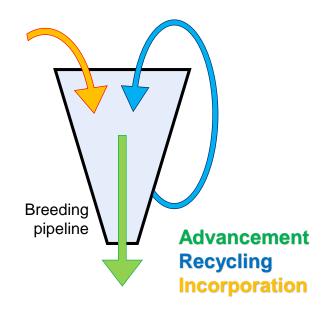




### Chasing the right signal

- Breeding value (GEBV)
  - Pattern: ADDITIVE GENETICS
  - Method: GBLUP, BayesABC, LASSO
  - Suits: RECYCLING, ADVANCEMENT

- Genomic value (EGV)
  - Pattern: ANY GENETICS
  - Method: RKHS, DNN, Random Forest
  - Suits: ADVANCEMENT





### Defining the problem: Metrics for success

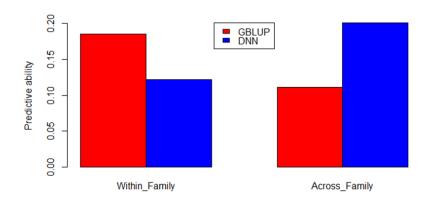
1. Scientist (why): to define the problem mathematically (easy to get it wrong)

2. Metric (what): Correlations, MSPE, Jaccard index, Accuracy, Success (1|0)

3. Testing (<a href="https://example.com/how">how</a>): Simulation or cross-validation (CV) on real data? WF vs AF?

How to design an adequate cross-validation???

### Two metrics, two different answers



Grain yield on the SoyNAM population: Populations 1:8 predicting populations 9:15

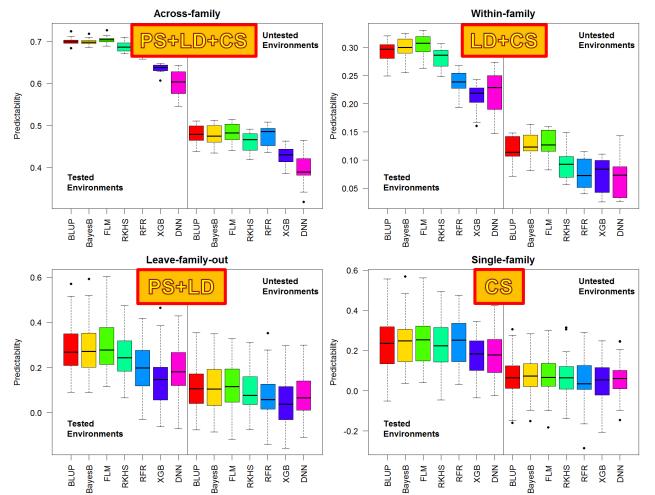


# Testing machines for different scenarios of genomic prediction

	Genotype	Environment	Prediction Difficulty
CV00	New	New	****
CV0	Observed	New	***
CV1	New	Observed	***
CV2	Observed	Observed	*

Adapted from Crossa et al. (2017) doi.org/10.1016/j.tplants.2017.08.011





SoyNAM data

ES: 2012 (7 loc)

PS: 2013 (4 loc)

#Fam = 40

Genos = 5600 SNPs = 4300

Obs: 3k-5k obs/loc



Untested

Tested

- Population structure (PS)
- Linkage disequilibrium (LD)
- Cosegregation / Haplotype (CS)



CV scheme

#### 1. Good

- Intro and motivation
- Filters in PB

#### 2. Bad

- Metrics of success
- Adequate validation

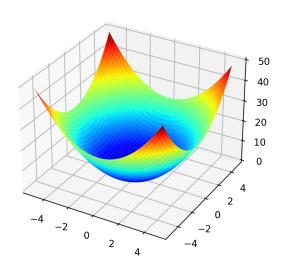
#### 3. Ugly

- Optimization
- From RR to NN

### Part 3 – Ugly learning

### Divergency in philosophy

- In quantitative genetics:
  - Parameters: Variances + Regression coefficients
  - Function: Likelihood (complex and convex)
  - Tuning: Generally not needed
  - Method: First order (EM, MCMC), second order (AI, MIVQUE)
- In machine learning:
  - Parameters: Regression coefficients
  - Function: MSE, L2 (simple)
  - Tuning: Cross validations, need <u>secondary objective function</u>
  - Method: First order: coordinate & gradient descent





### Solving: y = Xb + e

Finding  $\rightarrow$  argmin( $e'e + \lambda b'b$ )



I've created a monster!!

Coordinate descent

(Use diagonals of LHS)

$$b_{j}^{t+1} = \frac{x_{j}'(y - X_{-j}b_{-j})}{x_{j}'x_{j} + \lambda}$$

**Used for WGR (RR, BayesA)** 

glmnet, BGLR, bWGR, GS3

Gradient descent

(Does not build LHS)

$$b^{t+1} = b^t - \frac{2r}{n} [X'(y - Xb^t) + \lambda b^t]$$

Second order

(Builds entire LHS)

$$b = (X'X + \lambda)^{-1}(X'y)$$

**Used for Deep Neural Nets** 

TF/Keras, PyTorch, MXNet, h2o

Used for everything else

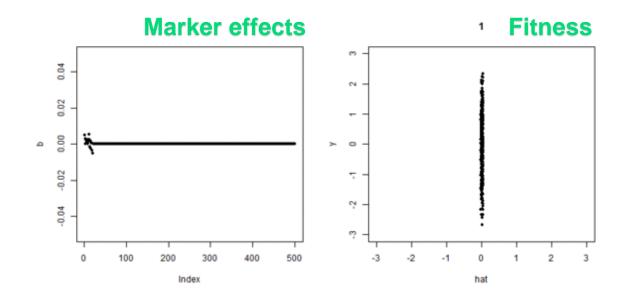
ASREML, Ime4, SAS, BLUPF90



### Coordinate descent of a RRBLUP



$$\begin{split} f(b_{RR}) &\to argmin \Big( \sum_{i=1}^{n} e^2 + \lambda \sum_{j=1}^{p} b^2 \Big) \\ SS &= e'e + \lambda(b'b) \\ SS &= (y - xb)'(y - xb) + \lambda(b'b) \\ SS &= y'y + (xb)'(xb) - 2(y'xb) + \lambda(b'b) \\ SS &= y'y + b'(x'x)b - 2(y'xb) + \lambda(b'b) \\ \frac{\partial SS}{\partial b} &= y'y + 2(x'x)b^2 - 2(y'x)b + 2\lambda b^2 \\ 0 &= 2b(x'x) - 2(y'x) + 2b\lambda \\ -2b(x'x + \lambda) &= -2(y'x) \\ b &= \frac{-2(y'x)}{-2(x'x + \lambda)} = \frac{y'x}{x'x + \lambda} \end{split}$$



### **Gradient descent of RRBLUP**

Scatter plot between true solution and solution at iteration x

#### **Small** learning rate

#### **Large** learning rate



$$y = X\beta + e$$
$$\beta^{t+1} = \beta^t - \alpha \nabla$$

• Where

$$\alpha$$
 = Learning rate

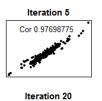
$$f = (y - Xb)'(y - Xb) + \lambda b'b$$

$$\nabla = f' = \frac{\partial f}{\partial b} = -2X'(y - Xb) + 2\lambda b$$

Thus

$$\nabla = n^{-1}(-2X'e + 2\lambda\beta^{t}) = -2n^{-1}(X'e - \lambda\beta^{t})$$

$$\beta^{t+1} = \beta^t + \frac{2\alpha}{r} (X'e - \lambda \beta^t)$$



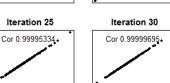
Iteration 35

Cor 0.99999991.



Iteration 40

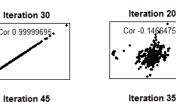
Cor 1

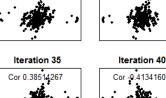


Cor 1

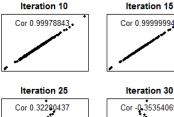
Iteration 15

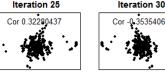
Cor 0.99781079 •

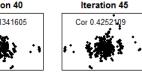




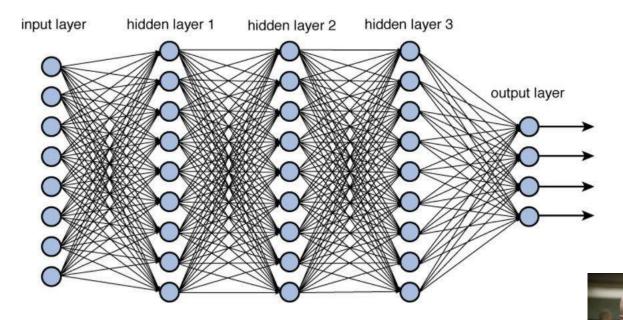
Iteration 5







#### **Deep Neural Network**



**WHAT IS** 

THIS NONSENSE?



### From Ridge Regression to Deep Neural Net

Models illustrated without intercept

Linear model

$$y = Xb + e$$

PLS/PCR model

$$y = (XB_1)b_2 + e$$

NN model

$$y = \alpha(XB_1)b_2 + e$$

Deep NN model

$$y = \alpha(\alpha(XB_1)B_2)b_3 + e$$

 $\alpha$  = activation function



### The Scary Deep Neural Network

$$Y = \alpha(\alpha(XB_1)B_2)B_3 + E$$

- Fit layers
  - $H_1 = \alpha(XB_1)$
  - $H_2 = \alpha(H_1B_2)$
  - $H_3 = H_2 B_3$
- Compute gradients (aka. get residuals)
  - $E_3 = Y H3$
  - $E_2 = \alpha(E_3 B_3')$
  - $E_1 = \alpha(E_2B_2')$
- Update coefficient
  - $B_1 = B_1 \gamma \left( \frac{2r_1}{n} \left[ X' E_1 \lambda B_1 \right] \right)$
  - $B_2 = B_2 \gamma \left( \frac{2r_2}{n} [H_1' E_2 \lambda B_2] \right)$
  - $B_3 = B_3 \gamma \left( \frac{2r_3}{n} [H_2' E_3 \lambda B_3] \right)$

#### Top-to-bottom code ~ 30 lines

```
dnn = function(y, X, nit=1000, batch=250, RELU=FALSE, Leak=0.1,
                     dropout=0, Lambda=0.1,LrnRate = 1,Nodes1=4, Nodes2=4){
   if(is.null(ncol(v))) v = matrix(v)
  muY = colMeans(y,na.rm=1); sdY = apply(y,2,sd,na.rm=1); y = apply(y,2,scale) ActFun = tanh; if(RELU) ActFun = function(x)\{x[x<0]=x[x<0]\} Leak; return(x)
   DropOut = function(x,prc=dropout) {x[sample(length(x),length(x)*prc)]:return(x)
   n = nrow(X); p = ncol(X); k = ncol(y)
  n1 = Nodes1; n2 = Nodes2; lmb = Lambda; rate = LrnRate/c(p,n1,n2)
 b1 = matrix(rnorm(n1*p,0,1/p),p,n1)
b2 = matrix(rnorm(n1*n2,0,1/n1),n1,n2)
b3 = matrix(rnorm(n2,0,1/n2),n2,k)
 CNV1 = CNV2 = CC
     w = sample(n,batch,replace=T); y0 = y[w,]; X0 = X[w,]
    Hl = ActFun(X0%%bl); H2 = ActFun(H1%%b2); H3 = H2%%b3
e3 = y0-H3; ff(anyMc(e3)) e3[fs.na(e3)] e2
e2 = ActFun(e3 %% t(b3)); e1 = ActFun(e2 %% t(b2))
     \begin{array}{lll} b1 = b1 + DropOut(t(X0)\%\%(e1) - lmb^*b1)^*(2/n)^*rate[\\ b2 = b2 + DropOut(t(H1)\%\%(e2) - lmb^*b2)^*(2/n)^*rate[\\ \end{array}
     b3 = b3 + DropOut(t(H2)%*%(e3))*(2/n)*rate[
     CNV1 = c(CNV1,mean(apply(e3,2,var,na.rm=
   CNV2 = c(CNV2,mean(diag(cor(H3,V0,use='p')))) out = list(af=ActFun,b1=b1,b2=b2,b3=b3,conv_MSE=CNV1,conv_GOF=CNV2,mu=muY,sd=sd)
predict.smalldnn = function(object,newdata)
   x = object; h = x$af(x$af(newdata%*%x$b1)%*%x$b2)%*%x$b3
   for(i in 1:ncol(h)) h[,i] = xsmu[i] + xssd[i]*h[,i]
```

### Thank you for your attention!

#### Remarks:

- ML can be a powerful tool for plant breeding
- 2) Cross-validation must be carefully design to understand the machinery
- 3) The scary black boxes are usually simple methods

### Questions??

Alencar Xavier

Alencar.Xavier@Corteva.com



### Some free lit to look up!

